

Database Systems

SQL: Data Manipulation (DML)

SQL Joins

SELECT Statement - Syntax

```

SELECT [DISTINCT | ALL]
  { * | [columnExpression [AS newName]]
    [, ...] }
FROM TableName [alias] [, ...]
[WHERE condition]
[GROUP BY columnList]
[HAVING condition]
[ORDER BY columnList]

```

- Syntax order of the clauses cannot be changed
- Only SELECT and FROM are mandatory

Keith Tang

COMP2714

2

SELECT Statement - Coding

Logical / Interpretation Order:

- FROM Specifies table(s) to be used
- WHERE Filters rows by some conditions
- GROUP BY Forms groups of rows with same column value
- HAVING Filters groups subject to some conditions
- SELECT Specifies which columns are to appear in output
- ORDER BY Specifies the order of the output

Keith Tang

COMP2714

3

Books Database Schema

authors (au_id, au_fname, au_lname, phone, address, city, state, zip)

title_authors (title_id, au_id, au_order, royalty_share)

publishers (pub_id, pub_name, city, state, country)

royalties (title_id, advance, royalty_rate)

titles (title_id, title_name, type, pub_id, pages, price, sales, pubdate, contract)

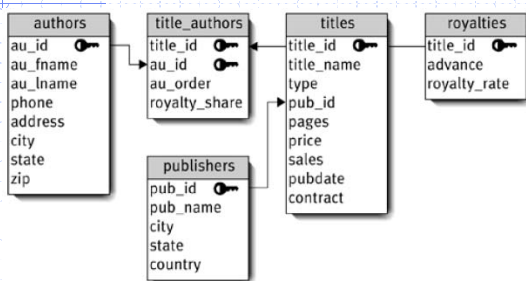
<http://www.fehily.com/books/SQL-Visual-QuickStart-Guide-3rd.html>

Keith Tang

COMP2714

4

Books Database Schema



Keith Tang

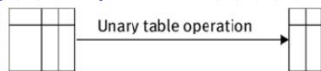
COMP2714

5

Relational Algebra

Relational algebraic operations:

- Projection
- Restriction



- Product
- Join
- Division
- Union
- Intersection
- Difference



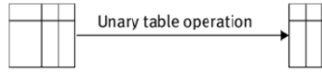
Keith Tang

COMP2714

6

Relational Algebra

◆ Unary operators:



- Projection – pick out columns
- SQL: SELECT <columns>
- Restriction – pick out rows meeting certain conditions
- SQL: WHERE <conditions>

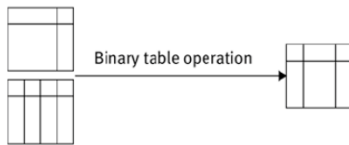
Keith Tang

COMP2714

7

Relational Algebra

◆ Binary operators



- Product – SQL: CROSS JOIN
- Join – SQL: INNER JOIN, OUTER JOIN
SQL: FROM clause
- Division – no single SQL counterpart (later)

Keith Tang

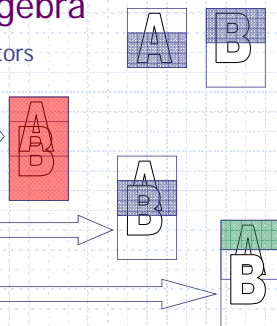
COMP2714

8

Relational Algebra

◆ Binary SET operators Union Compatible

- A Union B : All rows from both tables
- A Intersect B : Rows existed in both tables
- A Difference B : Rows in one but not in other table



Keith Tang

COMP2714

9

Cross Join / Product Example

```
SQL> SELECT * FROM mypubs;
```

```
PUB PUB_NAME
-----
P01 Pearson
P02 McGraw-Hill
P03 Morgan-Kaufmann
```

```
SQL> SELECT * FROM mytitles;
```

```
TIT TITLE_NAME  TYPE      PUB PUBLISHDATE
-----
T01 1977!       history   P01 2000-08-01
T02 World War I history   P03 1998-04-01
T03 Database    computer  P03 2000-09-01
T05 Your Ideas  psychology P01 2001-01-01
T10 Mr President biography P02
T13 Civil War   history   P02 1999-05-31
```

Keith Tang

COMP2714

10

SQL> SELECT * 2 FROM mypubs CROSS JOIN mytitles;				SQL> SELECT * 2 FROM mypubs, mytitles;			
PUB	PUB_NAME	TIT	TITLE_NAME	TYPE	PUB	PUBLISHDATE	
P01	Pearson	T01	1977!	history	P01	2000-08-01	
P01	Pearson	T02	World War I	history	P03	1998-04-01	
P01	Pearson	T03	Database	computer	P03	2000-09-01	
P01	Pearson	T05	Your Ideas	psychology	P01	2001-01-01	
P01	Pearson	T10	Mr President	biography	P02		
P01	Pearson	T13	Civil War	history	P02	1999-05-31	
P02	McGraw-Hill	T01	1977!	history	P01	2000-08-01	
P02	McGraw-Hill	T02	World War I	history	P03	1998-04-01	
P02	McGraw-Hill	T03	Database	computer	P03	2000-09-01	
P02	McGraw-Hill	T05	Your Ideas	psychology	P01	2001-01-01	
P02	McGraw-Hill	T10	Mr President	biography	P02		
P02	McGraw-Hill	T13	Civil War	history	P02	1999-05-31	
P03	Morgan-Kaufmann	T01	1977!	history	P01	2000-08-01	
P03	Morgan-Kaufmann	T02	World War I	history	P03	1998-04-01	
P03	Morgan-Kaufmann	T03	Database	computer	P03	2000-09-01	
P03	Morgan-Kaufmann	T05	Your Ideas	psychology	P01	2001-01-01	
P03	Morgan-Kaufmann	T10	Mr President	biography	P02		
P03	Morgan-Kaufmann	T13	Civil War	history	P02	1999-05-31	

18 rows selected.

COMP2714

11

Cross Join / Product

◆ SQL1:

```
■ SELECT [DISTINCT | ALL]
      { * | columnList }
FROM Table1, Table2;
```

◆ SQL2:

```
■ SELECT [DISTINCT | ALL]
      { * | columnList }
FROM Table1 CROSS JOIN Table2;
```

Keith Tang

COMP2714

12

Inner Join

- List the names of all publishers which have published a book:

```
SELECT *
FROM mypubs
INNER JOIN mytitles
ON mypubs.pub_id = mytitles.pub_id;
```

PUB	PUB_NAME	TIT	TITLE_NAME	TYPE	PUB	PUBDATE
P01	Pearson	T01	1977!	history	P01	2000-08-01
P03	Morgan-Kaufmann	T02	World War I	history	P03	1998-04-01
P03	Morgan-Kaufmann	T03	Database	computer	P03	2000-09-01
P01	Pearson	T05	Your Ideas	psychology	P01	2001-01-01
P02	McGraw-Hill	T10	Mr President	biography	P02	
P02	McGraw-Hill	T13	Civil War	history	P02	1999-05-31

6 rows selected.

Keith Tang

COMP2714

13

```
SQL> SELECT *
2 FROM mypubs CROSS JOIN mytitles
;
```

PUB	PUB_NAME	TIT	TITLE_NAME	TYPE	PUB	PUBDATE
P01	Pearson	T01	1977!	history	P01	2000-08-01
P01	Pearson	T02	World War I	history	P03	1998-04-01
P01	Pearson	T03	Database	computer	P03	2000-09-01
P01	Pearson	T05	Your Ideas	psychology	P01	2001-01-01
P01	Pearson	T10	Mr President	biography	P02	
P01	Pearson	T13	Civil War	history	P02	1999-05-31
P02	McGraw-Hill	T01	1977!	history	P01	2000-08-01
P02	McGraw-Hill	T02	World War I	history	P03	1998-04-01
P02	McGraw-Hill	T03	Database	computer	P03	2000-09-01
P02	McGraw-Hill	T05	Your Ideas	psychology	P01	2001-01-01
P02	McGraw-Hill	T10	Mr President	biography	P02	
P02	McGraw-Hill	T13	Civil War	history	P02	1999-05-31
P03	Morgan-Kaufmann	T01	1977!	history	P01	2000-08-01
P03	Morgan-Kaufmann	T02	World War I	history	P03	1998-04-01
P03	Morgan-Kaufmann	T03	Database	computer	P03	2000-09-01
P03	Morgan-Kaufmann	T05	Your Ideas	psychology	P01	2001-01-01
P03	Morgan-Kaufmann	T10	Mr President	biography	P02	
P03	Morgan-Kaufmann	T13	Civil War	history	P02	1999-05-31

Keith Tang

COMP2714

14

```
SQL> SELECT *
2 FROM mypubs CROSS JOIN mytitles;
SQL> SELECT *
2 FROM mypubs, mytitles;
```

PUB	PUB_NAME	TIT	TITLE_NAME	TYPE	PUB	PUBDATE
P01	Pearson	T01	1977!	history	P01	2000-08-01
P01	Pearson	T05	Your Ideas	psychology	P01	2001-01-01
P02	McGraw-Hill	T10	Mr President	biography	P02	
P02	McGraw-Hill	T13	Civil War	history	P02	1999-05-31
P03	Morgan-Kaufmann	T02	World War I	history	P03	1998-04-01
P03	Morgan-Kaufmann	T03	Database	computer	P03	2000-09-01

6 rows selected.

Keith Tang

COMP2714

15

Inner Join

- ◆ List the names of all publishers which have published a history book:

```
◆ SELECT mypubs.pub_id, pub_name
   FROM mypubs
      INNER JOIN mytitles
        ON mypubs.pub_id = mytitles.pub_id
  WHERE type = 'history';
```

```
PUB  PUB_NAME
---  -
P01  Pearson
P03  Morgan-Kaufmann
P02  McGraw-Hill
```

Keith Tang

COMP2714

16

Qualifying Column Names

- ◆ Need to qualify column names when there is ambiguity
- ◆ Can use table names, except in self-join
- ◆ Can use an alias for a table named in the FROM clause, separated from table name with a space; AS keyword is optional
- ◆ Cannot mix – once aliases are defined, must use aliases to qualify column names

Keith Tang

COMP2714

17

Inner Join Using Table Aliases

- ◆ List the names of all publishers which have published a history book:

```
◆ SELECT p.pub_id, pub_name
   FROM mypubs p
      INNER JOIN mytitles t
        ON p.pub_id = t.pub_id
  WHERE type = 'history';
```

```
PUB  PUB_NAME
---  -
P01  Pearson
P03  Morgan-Kaufmann
P02  McGraw-Hill
```

Keith Tang

COMP2714

18

Types of Join

- ◆ CROSS JOIN (no conditions)
- ◆ [INNER] JOIN
 - EQUI-JOIN (ON conditions are equalities)
 - THETA-JOIN (any comparison ON conditions)
- ◆ NATURAL JOIN (equality of common columns)
- ◆ SELF-JOIN (mostly to other rows of same table)
- ◆ [OUTER] JOIN
 - LEFT [OUTER] JOIN
 - RIGHT [OUTER] JOIN
 - FULL [OUTER] JOIN

Keith Tang

COMP2714

19

INNER JOIN vs NATURAL JOIN

- ◆ SQL2 provides alternative ways for inner joins:
 1. `FROM mypubs p JOIN mytitles t ON p.pub_id = t.pub_id`
 2. `FROM mypubs p JOIN mytitles t USING (pub_id)`
 3. `FROM mypubs NATURAL JOIN mytitles`
- ◆ Note1: uses FROM clause only, not WHERE clause
- ◆ Note2: 1st format produces result with 2 pub_id columns
- ◆ Note3: 2nd and 3rd formats gives result with only 1 pub_id column (eliminates duplicates)

Keith Tang

COMP2714

20

Which JOIN Syntax?

- ◆ If result columns come from more than 1 table, must have a way to combine them.
- ◆ Older SQL-89
 - FROM : tables separated by commas
 - WHERE : the join column(s) comparisons
- ◆ Newer SQL-92
 - FROM : tables and join column(s)
 - There are multiple syntax alternatives

Keith Tang

COMP2714

21

Which JOIN Syntax?

◆ Older SQL-89

```
SELECT au_fname, au_lname, a.city
  FROM authors a, publishers p
 WHERE a.city = p.city;
```

◆ Newer SQL-92

```
SELECT au_fname, au_lname, a.city
  FROM authors a
 INNER JOIN publishers p
    ON a.city = p.city;
```

Keith Tang

COMP2714

22

Sorting a Join Result

- ◆ List the names of all publishers which have published a history book:

```
◆ SELECT mypubs.pub_id, pub_name
   FROM mypubs
      INNER JOIN mytitles
        ON mypubs.pub_id = mytitles.pub_id
  WHERE type = 'history'
  ORDER BY pub_name;
```

```
PUB PUB_NAME
---
P02 McGraw-Hill
P03 Morgan-Kaufmann
P01 Pearson
```

Keith Tang

COMP2714

23

Three Table Join – SQL-89

```
SELECT t.title_id, t.pub_id,
       p.pub_name, r.advance
  FROM publishers p, titles t,
       royalties r
 WHERE p.pub_id = t.pub_id
    AND t.title_id = r.title_id
    AND r.advance < 20000;
```

Keith Tang

COMP2714

24

Three Table Join – SQL-92

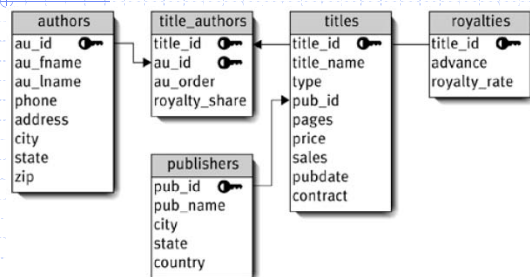
```
SELECT t.title_id, t.pub_id,  
       p.pub_name, r.advance  
FROM publishers p  
      INNER JOIN titles t  
        ON p.pub_id = t.pub_id  
      INNER JOIN royalties r  
        ON t.title_id = r.title_id  
WHERE r.advance < 20000;
```

Keith Tang

COMP2714

25

Books Database Schema



Keith Tang

COMP2714

26

Computing an Inner Join

◆ Procedure for generating the result of a join CONCEPTUALLY:

1. Form CROSS JOIN of the tables named in FROM clause
2. If WHERE clause, apply the conditions to each row, retaining rows meeting conditions
3. Apply SELECT list to retain those columns
4. If DISTINCT, eliminate any duplicate rows
5. If ORDER BY clause, sort result table as required

Keith Tang

COMP2714

27

A JOIN Example

- ◆ Give me a list of the authors and publishers involved in publishing computer books.

authors (au_id, au_fname, au_lname, phone, address, city, state, zip)

title_authors (title_id, au_id, au_order, royalty_share)

publishers (pub_id, pub_name, city, state, country)

royalties (title_id, advance, royalty_rate)

titles (title_id, title_name, type, pub_id, pages, price, sales, pubdate, contract)

Keith Tang

COMP2714

28

A JOIN Example

- ◆ Give me a list of the authors and publishers involved in publishing computer books.

Keith Tang

COMP2714

29

A JOIN Example

- ◆ Give me a list of the authors and publishers involved in publishing computer books.

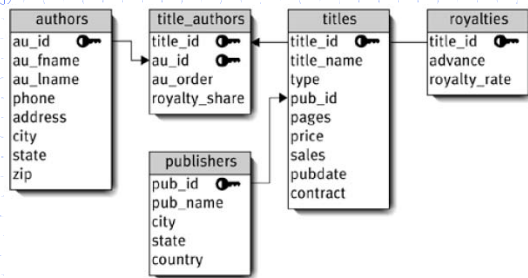
1. FROM ??? -- 3 questions

Keith Tang

COMP2714

30

Give me a list of the authors and publishers involved in publishing computer books.



Keith Tang

COMP2714

31

Outer Joins

- ◆ If one row of a joined table is unmatched, that row is omitted from result table
- ◆ Outer join operations retain rows that do not satisfy the join conditions
- ◆ Consider the following query:
 - Find the publishers who have not yet published a book.
 - Find the books without a publisher lined up as yet

Keith Tang

COMP2714

32

Find the publishers who have not yet published a book

```
SQL> SELECT * FROM mypubs;
```

```
PUB PUB_NAME
-----
P01 Pearson
P02 McGraw-Hill
P03 Morgan-Kaufmann
P04 New World Book
```

```
SQL> SELECT * FROM mytitles;
```

```
TIT TITLE_NAME  TYPE      PUB PUBDATE
-----
T01 1977!       history   P01 2000-08-01
T02 World War I history   P03 1998-04-01
T03 Database    computer  P03 2000-09-01
T05 Your Ideas  psychology P01 2001-01-01
T10 Mr President biography
T13 Civil War   history   P02 1999-05-31
```

Keith Tang

COMP2714

33

Find the publishers who have not yet published a book

◆ An INNER JOIN can never help to answer such queries:

```
◆ SELECT *
  FROM mypubs p
    JOIN mytitles t      -- INNER JOIN
      ON p.pub_id = t.pub_id;
```

PUB	PUB_NAME	TIT	TITLE_NAME	TYPE	PUB	PUBDATE
P01	Pearson	T01	19771	history	P01	2000-08-01
P03	Morgan-Kaufmann	T02	World War I	history	P03	1998-04-01
P03	Morgan-Kaufmann	T03	Database	computer	P03	2000-09-01
P01	Pearson	T05	Your Ideas	psychology	P01	2001-01-01
P02	McGraw-Hill	T13	Civil War	history	P02	1999-05-31

Keith Tang

COMP2714

34

LEFT [OUTER] JOIN Example

◆ Find the publishers who have not yet published a book (& what is the percentage of publishers)

```
◆ SELECT *
  FROM mypubs p
    LEFT JOIN mytitles t      -- LEFT JOIN
      ON p.pub_id = t.pub_id
 WHERE t.pub_id IS NULL;
```

PUB	PUB_NAME	TIT	TITLE_NAME	TYPE	PUB	PUBDATE
P01	Pearson	T01	19771	history	P01	2000-08-01
P03	Morgan-Kaufmann	T02	World War I	history	P03	1998-04-01
P03	Morgan-Kaufmann	T03	Database	computer	P03	2000-09-01
P01	Pearson	T05	Your Ideas	psychology	P01	2001-01-01
P02	McGraw-Hill	T13	Civil War	history	P02	1999-05-31
P04	New World Book					

Keith Tang

COMP2714

35

Outer Join Result Includes:

- ◆ INNER JOIN result rows
(can have multiple matched rows)
- ◆ Additional unmatched rows with NULL-fills
(only 1 row per unmatched row)
 - LEFT JOIN – extra rows from LEFT table
 - RIGHT JOIN – extra rows from RIGHT table
 - FULL JOIN – extra rows from both tables

Keith Tang

COMP2714

36

RIGHT [OUTER] JOIN Example

- ◆ Find the books without a publisher lined up as yet

```
◆ SELECT *
  FROM mypubs p
    RIGHT JOIN mytitles t
      ON p.pub_id = t.pub_id
 WHERE p.pub_id IS NULL;
```

PUB	PUB_NAME	TIT	TITLE_NAME	TYPE	PUB	PUBDATE
P01	Pearson	T01	1977!	history	P01	2000-08-01
P03	Morgan-Kaufmann	T02	World War I	history	P03	1998-04-01
P03	Morgan-Kaufmann	T03	Database	computer	P03	2000-09-01
P01	Pearson	T05	Your Ideas	psychology	P01	2001-01-01
P02	McGraw-Hill	T13	Civil War	history	P02	1999-05-31
		T10	Mr President biography			

Keith Tang

COMP2714

37

FULL [OUTER] JOIN Example

- ◆ I can match the publishers who have not yet published a book with books without a publisher

```
◆ SELECT *
  FROM mypubs p
    FULL JOIN mytitles t
      ON p.pub_id = t.pub_id
 WHERE ; -- what ???
```

PUB	PUB_NAME	TIT	TITLE_NAME	TYPE	PUB	PUBDATE
P01	Pearson	T01	1977!	history	P01	2000-08-01
P03	Morgan-Kaufmann	T02	World War I	history	P03	1998-04-01
P03	Morgan-Kaufmann	T03	Database	computer	P03	2000-09-01
P01	Pearson	T05	Your Ideas	psychology	P01	2001-01-01
P02	McGraw-Hill	T13	Civil War	history	P02	1999-05-31
P04	New World Book					
		T10	Mr President biography			

Keith Tang

COMP2714

38

Self Join

emp_id	emp_name	boss_id
E01	Lord Copper	NULL
E02	Jocelyn Hitchcock	E01
E03	Mr. Salter	E01
E04	William Boot	E03
E05	Mr. Corker	E03

- ◆ List each employee with his/her boss

```
◆ SELECT e1.emp_name AS "Employee",
       e2.emp_name AS "Boss"
       stf.staffNo, stf.fName, stf.lName
  FROM employee e1
    JOIN employee e2
      ON e1.boss_id = e2.emp_id;
```

Keith Tang

COMP2714

39

Self Join

- ◆ List each employee with his/her boss

```
◆ SELECT e1.emp_name AS "Employee",  
        e2.emp_name AS "Boss"  
        stf.staffNo, stf.fName, stf.lName  
FROM employee e1  
JOIN employee e2  
ON e1.boss_id = e2.emp_id;
```

Employee	Boss
Jocelyn Hitchcock	Lord Copper
Mr. Salter	Lord Copper
William Boot	Mr. Salter
Mr. Corker	Mr. Salter

Keith Tang

COMP2714

40

Self Join

- ◆ List each employee with his/her boss

```
◆ SELECT e1.emp_name AS "Employee",  
        e2.emp_name AS "Boss"  
        stf.staffNo, stf.fName, stf.lName  
FROM employee e1  
LEFT JOIN employee e2  
ON e1.boss_id = e2.emp_id;
```

Employee	Boss
Lord Copper	
Jocelyn Hitchcock	Lord Copper
Mr. Salter	Lord Copper
William Boot	Mr. Salter
Mr. Corker	Mr. Salter

Keith Tang

COMP2714

41